

Phase 1: Foundations and Classic ML (Weeks 1–4)

Week	Focus Area	Key Concepts	Recommended Resources (Free/Low-Cost)	Project / Practice Goal
1	Python & Tools	Python Fundamentals (OOP, functions), NumPy (arrays, matrix operations), Environment Setup (Jupyter/Anaconda).	W3Schools Python Tutorial, NumPy Documentation, Kaggle's Python track .	Practice matrix math with NumPy; create a simple Python class.
2	Data Analysis & Visualization	Pandas (DataFrames, cleaning, manipulation), Matplotlib & Seaborn (Basic plots: scatter, histogram), Git/GitHub basics.	Kaggle's Pandas/Data Viz tracks , freeCodeCamp's Data Analysis Course.	Complete an Exploratory Data Analysis (EDA) on a simple dataset (e.g., Titanic).
3	Math Essentials	Linear Algebra (vectors, matrices), Calculus (introduction to Gradient Descent), Probability & Statistics (mean, variance, distributions).	3Blue1Brown's Essence of Linear Algebra , Khan Academy (Statistics/Probability), StatQuest (YouTube) .	Solve basic statistical problems and plot probability distributions in Python.

4	ML Basics & Regression	Introduction to ML (Supervised vs. Unsupervised), Train/Test Split, Linear & Logistic Regression , Evaluation Metrics (MSE, R^2 , Accuracy).	Andrew Ng's Machine Learning Specialization (Coursera) Course 1, Google's ML Crash Course.	Build a Linear Regression model to predict a simple variable (e.g., house price).
---	-----------------------------------	---	---	--

Phase 2: Intermediate Algorithms & Deep Learning (Weeks 5–8)

Week	Focus Area	Key Concepts	Recommended Resources (Free/Low-Cost)	Project / Practice Goal
5	Classification & Model Tuning	Decision Trees, Random Forests, SVM , Evaluation Metrics (Precision, Recall, F1-Score), Cross-Validation & Hyperparameter Tuning .	Scikit-learn documentation/user guide, Kaggle's Intermediate ML Course .	Build a Random Forest classifier for the Titanic or a similar classification dataset.
6	Unsupervised Learning & Clustering	K-Means Clustering , Hierarchical Clustering, Dimensionality Reduction (PCA) , Association Rules.	Scikit-learn Unsupervised Learning modules.	Perform Customer Segmentation using K-Means and use PCA to visualize the results.

Week	Focus Area	Key Concepts	Recommended Resources (Free/Low-Cost)	Project / Practice Goal
7	Deep Learning Foundations	The Perceptron, Neural Networks (NNs) architecture, Activation Functions (Sigmoid, ReLU), TensorFlow/Keras setup.	DeepLearning.AI's Deep Learning Specialization (Course 1), TensorFlow Tutorials .	Build a simple Feedforward NN to classify the MNIST handwritten digits dataset.
8	Convolutional Neural Networks (CNNs)	Convolution/Pooling layers, Basic CNN Architectures, Transfer Learning fundamentals for images.	fast.ai Practical Deep Learning for Coders , Specific TensorFlow/PyTorch CNN tutorials .	Implement an Image Classifier or utilize a pre-trained model for a small image dataset.

Phase 3: Advanced Topics and Deployment (Weeks 9–12)

Week	Focus Area	Key Concepts	Recommended Resources (Free/Low-Cost)	Project / Practice Goal
9	Recurrent Networks & Sequences	Recurrent Neural Networks (RNNs) , LSTMs (Long Short-Term Memory) and	DeepLearning.AI's Specialization (Sequence Models), Tutorials on time series prediction.	Build a simple Time Series Predictor or a basic Text Generator using LSTMs.

		GRUs, handling sequence data.		
10	Natural Language Processing (NLP)	Text Preprocessing, Tokenization, Word Embeddings (Word2Vec, GloVe) , Basic Transformer architecture.	Andrew Ng's NLP Specialization , Hugging Face Introduction to Transformers.	Build a Sentiment Analyzer using an LSTM model on a text review dataset.
11	MLOps & Deployment	Model saving (Pickle), API creation (Flask/FastAPI) , Introduction to Docker (containerization).	Tutorials for deploying ML models with Flask/Streamlit, Practical MLOps guides.	Deploy your Sentiment Analyzer model as a simple web API using Flask or FastAPI.
12	Final End-to-End Project	Integrating all skills: Data acquisition, pre-processing, modeling, hyperparameter tuning, documentation, and presentation.	Kaggle Competitions , GitHub for portfolio management.	Build a complete, fully documented project (e.g., a personalized recommendation engine) and make it public on GitHub.

What's Next? Your Path After 12 Weeks

Career Action	Specific Goal / Outcome	Tools & Resources
Deepen Portfolio	Build 3-5 Polished Projects focused on a single domain (e.g., Computer Vision, Advanced NLP). Ensure at least one project is fully deployed and live.	Kaggle Competitions, GitHub, Streamlit / Hugging Face Spaces for easy deployment, MLflow for experiment tracking.
Advanced Specialization	Master a specific advanced track based on your interest: MLOps (production code) or Deep Generative Models (research/algorithms).	Tutorials on Kubernetes/Docker , MLOps principles, arXiv (for reading research papers), Advanced Andrew Ng Specializations .
Launch Job Search	Integrate your new projects into your resume, perfect your tech interview skills, and start networking efficiently.	UpStride Cohort 8 Materials (Resume/Branding), LeetCode/HackerRank (DSA practice), UpStride Mentor (for mock system design interviews and introductions).
Continuous Learning	Stay current with the rapidly evolving AI landscape and continuously iterate on your skills.	ArXiv (for new papers), Key AI Newsletters , DeepLearning.AI Short Courses , Industry blogs (e.g., Google AI, Meta AI).